

LECTURE 11 SQL BASICS

Introduction

- SQL is the standard language for querying and manipulating (relational) data
(<https://www.sql50.com>)
 - SQL stands for Structured Query Language
 - Initially developed at IBM by Donald Chamberlin and Raymond Boyce in the early 1970s, and called **SEQUEL** (*Structured English Query Language*)
 - Many standards out there: SQL92, SQL3, SQL99,, **SQL2016**
 - (**Caveats**) Vendors support various subsets of these standards
- Why SQL?
 1. A **very-high-level, declarative language**, in which the programmer is able to avoid specifying a *lot of data-manipulation details* that would be necessary in languages like C++
 2. Its queries are “**optimized**” quite well over the years, yielding efficient query execution and performance

Components of SQL

- DDL – Data Definition Language
 - Define and modify schema
- DML – Data Manipulation Language
 - Query (glorified search) data
 - Insert data
 - Remove data
 - Update data

Some SQL syntax

- SQL is *case-insensitive*
 - In general, upper- and lower-case characters are the same, except inside quoted strings
- Single quote for a string. Two single quotes inside a string represent the single-quote (apostrophe)
- Each statement finishes with a semicolon (;).

Create Table

Format: CREATE TABLE <name> (
 <list of elements>
);

- Each element is a pair consisting of an attribute and a type
- The most common types are:
 - INT or INTEGER (synonyms)
 - REAL or FLOAT (synonyms)
 - CHAR(n) = fixed-length string of n characters
 - VARCHAR(n) = variable-length string of up to n characters

Create Table - example

```
CREATE TABLE Sells (
    bar      CHAR(20) ,
    beer     VARCHAR(20) ,
    price    REAL
) ;
```

Declaring keys

- An attribute or list of attributes may be declared PRIMARY KEY or UNIQUE
 - Each says the attribute(s) so declared functionally determines all the attributes of the relation schema
 - There can be **only one PRIMARY KEY** for a relation, but **several UNIQUE attributes**
 - No attribute of a PRIMARY KEY can ever be **NULL** in any tuple. But attributes declared UNIQUE may have NULL's, and there may be several tuples with NULL
 - Single attribute keys

```
CREATE TABLE Beers (
    name CHAR(20) UNIQUE,
    manf CHAR(20)
);
```

Multi-attribute keys

```
CREATE TABLE Sells (
    bar        CHAR(20) ,
    beer       VARCHAR(20) ,
    price      REAL,
    PRIMARY KEY (bar, beer)
) ;
```

Other Declarations for Attributes

- Two other declarations we can make for an attribute are:
 - **NOT NULL** means that the value for this attribute may never be NULL
 - **DEFAULT <value>** says that if there is no specific value known for this attribute's component in some tuple, use the stated <value>

```
CREATE TABLE Drinkers (
    name CHAR(30) PRIMARY KEY,
    addr CHAR(50)
        DEFAULT '123 Monroe St.',
    phone CHAR(16)
);
```

Other useful SQL DDL commands

- Add an attribute to a table

```
ALTER TABLE Bars ADD  
    phone CHAR(10) DEFAULT 'unlisted' ;
```

- Remove an attribute from a table

```
ALTER TABLE Bars  
    DROP license;
```

- Remove a table

```
DROP TABLE Bars
```

Table manipulation commands summary

- Create a table
- Remove a table
- Modify a table:
 - Add fields to a table
 - Remove fields from a table

Insert a new record into a table

- The insert into statement inserts a new record into a table
 - `INSERT INTO table_name (column1, column2, column3, ...)`
`VALUES (value1, value2, value3, ...);`
 - `INSERT INTO table_name`
`VALUES (value1, value2, value3, ...);`

Insert into example

```
CREATE TABLE Drinkers (
    name CHAR(30) PRIMARY KEY,
    addr CHAR(50)
        DEFAULT '123 Monroe St.',
    phone CHAR(16)
);

Insert into Drinkers(name, addr, phone)
Values ('John Smith', '201 Marty Ln', '555-666-7777')
```

NAME	ADDR	PHONE
John Smith	201 Marty Ln	555-666-7777

Insert into example

- o Insert into Drinkers(name, phone)
Values ('Sophie Green', '111-222-3333')
- o Insert into Drinker(name)
Values ('John Doe')

NAME	ADDR	PHONE
John Smith	201 Marty Ln	555-666-7777
Sophie Green	123 Monroe St.	111-222-3333
John Doe	123 Monroe St.	NULL

Write SQL commands to create the following table

Employees	EmployeeID	LastName	FirstName	DeptID	Age
	001	John	Smith	100	20
	002	Why	Me	NULL	NULL
	003	What	NULL	NULL	18

SQL query

- SELECT-FROM-WHERE statements

- SELECT: desired attributes, expressions from the attributes
- FROM: one or more tables
- WHERE: condition about tuples of the table
- Using table Beers (name, manf) to find all beers made by Busch.

```
SELECT name  
FROM Beers  
WHERE manf = 'Busch'
```

- The answer is a relation with a single attribute name, and tuples with the name of each beer by Busch, such as Bud

Single-Relation Query

- Operations

1. Begin with the relation in the **FROM** clause
2. Apply the condition indicated by the **WHERE** clause to each tuple
3. Select the fields (or expressions from the fields) indicated by the **SELECT** clause

- Semantics

1. To implement this algorithm, think of a **tuple variable** ranging over each tuple of the relation mentioned in FROM
2. Check if the “current” tuple satisfies the WHERE clause
3. If so, compute the attributes or expressions of the SELECT clause using the components of this tuple

Single-relation query example

Beers

Name	manf
'Bud'	'Busch'
'Miller Lite'	'MillerCoors'
'Bud Lite'	'Busch'
'Blue Moon'	'Molson Coors'
'Corona Extra'	'Constellation Brands AB InBev'
'Michelob'	'Busch'

```
SELECT name  
FROM Beers  
WHERE manf = 'Busch'
```



Name	manf
'Bud'	'Busch'
'Bud Lite'	'Busch'
'Michelob'	'Busch'



Name
'Bud'
'Bud Lite'
'Michelob'

* In SELECT clauses

- When there is one relation in the FROM clause, * in the SELECT clause stands for “all attributes of this relation.”
- Example using Beers(name, manf):

```
SELECT *
FROM Beers
WHERE manf = 'Busch' ;
```

Name	manf
'Bud'	'Busch'
'Bud Lite'	'Busch'
'Michelob'	'Busch'

Renaming Attributes

- If you want the result to have different attribute names, use “AS <new name>” to rename an attribute
- Example using Beers(name, manf):

```
SELECT name AS beer, manf  
FROM Beers  
WHERE manf = 'Busch';
```

Beer	manf
'Bud'	'Busch'
'Bud Lite'	'Busch'
'Michelob'	'Busch'

Expressions in SELECT clauses

- Any expression that makes sense can appear as an element of a SELECT clause
- Example: from Sells(bar, beer, price):

```
SELECT bar, beer, price * 132 AS priceInYen  
FROM Sells;
```

Sells

Bar	Beer	Price
'Sue Bar'	'Bud'	1.00
'Joe Bar'	'Bud'	2.00
'Fat Daddy'	'Bud'	3.00



Bar	Beer	PricelnYen
'Sue Bar'	'Bud'	132.00
'Joe Bar'	'Bud'	264.00
'Fat Daddy'	'Bud'	396.00

Complex Conditions in WHERE Clause

- What you can use in WHERE:

- attribute names of the relation(s) in FROM
- comparison operators: $=, <>, <, >, \leq, \geq$
- apply arithmetic operations: stockprice \ast 2
- operations on strings (e.g., “ \mid ” for concatenation)
- Lexicographic order on strings
- Pattern matching: $s \text{ LIKE } p$
- Special stuff for comparing dates and times.
- Conditions in the WHERE clause can use AND, OR, NOT, and parentheses in the usual way
Boolean conditions are built

Complex Conditions in WHERE Clause

- From Sells(bar, beer, price), find the price Joe's Bar charges for “cheap” beers:

```
SELECT price  
FROM Sells  
WHERE bar = 'Joe Bar' AND  
      price < 5.0;
```

Pattern in WHERE clause

- WHERE clauses can have conditions in which a string is compared with a pattern, to see if it matches
- General form: <Attribute> LIKE <pattern> or <Attribute> NOT LIKE <pattern>
 - Pattern is a quoted string with % = “any string”; _ = “any character.”
 - ❖ From Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name  
FROM Drinkers  
WHERE phone LIKE '%555-__ __';
```

NULL value

- Tuples in SQL relations can have NULL as a value for one or more components. Its Meaning depends on context; Two common cases:
 - **Missing value** : e.g., we know Joe's Bar has some address, but we don't know what it is
 - **Inapplicable** : e.g., the value of attribute spouse for an unmarried person
- When any value is compared with NULL, the result value is an **UNKNOWN**
- The logic of conditions in SQL is really a 3-valued logic: TRUE, FALSE, UNKNOWN
 - A query only produces a tuple in the answer if its value for the WHERE clause is **TRUE** (not FALSE or UNKNOWN)

Example

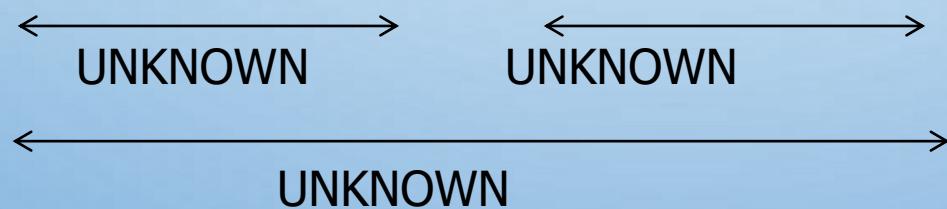
Sells

bar	beer	Price
Joe's	Bud	NULL

SELECT BAR

FROM SELLS

WHERE PRICE < 2.00 OR PRICE >= 2.00



Three-valued logic

- To understand how AND, OR, and NOT work in 3-valued logic, think of TRUE = 1, FALSE = 0, and UNKNOWN = $\frac{1}{2}$, AND = MIN; OR = MAX, NOT(x) = $1-x$.
- Example:

TRUE AND (FALSE OR NOT(UNKNOWN))

$$= \text{MIN}(1, \text{MAX}(0, (1 - \frac{1}{2})))$$

$$= \text{MIN}(1, \text{MAX}(0, \frac{1}{2}))$$

$$= \text{MIN}(1, \frac{1}{2})$$

$$= \frac{1}{2}$$

Example

Sells (bar, beer, price)

bar	beer	Price
Joe's	Bud Lite	NULL
Joe's	Bud	NULL
Big Mama	Bud	2.00

```
SELECT bar  
FROM Sells  
WHERE beer = 'Bud' OR price < 3.00
```

Explicit test for NULL

- Can test for NULL explicitly:

- o $x \text{ IS NULL}$
- o $x \text{ IS NOT NULL}$

```
SELECT *
FROM Person
WHERE age < 25 OR age >= 25 OR age IS NULL
```

Exercise

- From Sells(bar, beer, price), find the bars that sell ‘Bud Lite’ for less than 4 dollars.

Exercise

- From Students(Sid, SSN, LName, FName, Phone, GPA), find the student id, last name, first name, and GPA of all students whose GPA is above 3.5 and whose phone number has an area code of 850 (phone number format: xxx-xxx-xxxx).

Multi-relation Queries

- Interesting queries often combine data from more than one relation, we can address several relations in one query by listing them all in the FROM clause.
- Basic multi-relation queries:
 - The Cartesian product of the multiple relations in the FROM clause:
 - Use relations **Likes**(drinker, beer) and **Frequents**(drinker, bar) for the following query. Let **Likes** have 4 tuples and **Frequents** have 6 tuples, how many columns are in the output of the following query? How many tuples?

```
SELECT *
FROM Likes, Frequents
```

Multi-relation Queries

- Attributes in different relations may have the same name. Distinguish attributes of the same name by “<relation>.<attribute>”
- Using relations **Likes**(drinker, beer) and **Frequents**(drinker, bar) for the following query. Let **Likes** have 4 tuples and **Frequents** have 6 tuples, how many columns are in the output of the following query? How many tuples?

```
SELECT Likes.drinker, Likes.beer  
FROM Likes, Frequents
```

Multi-relation Queries

- **Query:** Using relations **Likes(drinker, beer)** and **Frequents(drinker, bar)**, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT Likes.beer  
FROM Likes, Frequents  
WHERE Frequents.bar = 'Joe''s' AND  
Frequents.drinker = Likes.drinker;
```

Semantics

- Almost the same as for single-relation queries:
 1. Start with **the (Cartesian) product** of all the relations in the FROM clause
 2. Apply the **selection** condition from the WHERE clause
 3. **Project** onto the list of attributes and expressions in the SELECT clause

```
SELECT a1, a2, ..., ak  
FROM   R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE  Conditions
```

Semantics

```
SELECT a1, a2, ..., ak  
FROM   R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE  Conditions
```

```
Answer = {}  
for x1 in R1 do  
  for x2 in R2 do  
    ....  
    for xn in Rn do  
      if Conditions  
        then Answer = Answer U {(a1,...,ak)}  
return Answer
```

Multi-relation queries exercise

- **Query:** Using relations **Likes(drinker, beer)** and **Sells(bar, beer, price)**, find the bars that sell beers that Sally likes.

Explicit Tuple Variables

- Sometimes, a query needs to use two copies of the same relation
 - Distinguish copies by following the relation name with the name of a tuple-variable, in the FROM clause
 - It's always an option to rename relations this way, even when not essential

```
SELECT s1.bar, s1.beer, s1.price  
FROM Sells s1, Sells s2  
WHERE s1.beer = s2.beer AND  
      s1.price < s2.price;
```

Exercise

- From students(Sid, SSN, LName, FName, Phone, GPA) and takes(Sid, course), what does the following query do?

```
SELECT t.course  
FROM students s, takes t  
WHERE s.FName = 'Ashley' AND s.LName = 'Brown' AND  
s.sid = t.sid;
```

SubQueries

- A parenthesized SELECT-FROM-WHERE statement (**subquery**) can be used as a value in a number of places, including FROM and WHERE clauses
 - Example: in place of a relation in the FROM clause, we can place another query, and then query its result
 - ❖ Better use a tuple-variable to name tuples of the result
- Subqueries that return Scalar
 - If a subquery is guaranteed to produce **one tuple with one component**, then the subquery can be used as a value
 - ❖ “Single” tuple often guaranteed by the key constraint
 - ❖ A run-time error occurs if there is no tuple or more than one tuple

Subquery example

- From Sells(bar, beer, price), find the bars that serve Miller for the same price Joe charges for Bud
 1. Find the price Joe charges for Bud
 2. Find the bars that serve Miller at that price

```
SELECT bar
FROM Sells
WHERE beer = 'Miller' AND
      price = (SELECT price
                FROM Sells
               WHERE bar = 'Joe Bar'
                 AND beer = 'Bud')
```

The IN Operator

- <tuple> IN <relation> is true if and only if the tuple is a member of the relation
 - <tuple> NOT IN <relation> means the opposite
 - IN-expressions can appear in WHERE clauses
 - The <relation> is often a subquery

Query: From Beers(name, manf) and Likes(drinker, beer), find the name and manufacturer of each beer that Sally likes

```
SELECT *
  FROM Beers
 WHERE name IN (      SELECT beer
                      FROM Likes
                     WHERE drinker = 'Sally'
                ) ;
```

Exercise

- **Query:** Using relations **Likes(drinker, beer)** and **Frequents(drinker, bar)**, **Sells(bar, beer, price)**, find the beers that some Joe's frequent customers like, but not sold by Joe's.

Exercise

- **Query:** Using relations **Likes(drinker, beer)** and **Frequents(drinker, bar)**, **Sells(bar, beer, price)**, find the beers that some Joe's frequent customers like, but not sold by Joe's.

```
SELECT Likes.beer
FROM Likes, Frequents
WHERE Likes.drinker = Frequents.drinker AND
frequents.bar = 'Joe''s'
AND likes.beer NOT IN (SELECT beer
FROM Sells
WHERE bar = 'Joe''s');
```

The EXISTS Operator

- EXISTS(<relation>) is true if and only if the <relation> is not empty
 - Being a Boolean-valued operator, EXISTS can appear in WHERE clauses

Query: From Beers(name, manf), find those beers that are the only beer by their manufacturer

```
SELECT name  
  FROM Beers b1  
 WHERE NOT EXISTS (  
       SELECT *  
         FROM Beers  
        WHERE manf = b1.manf AND  
              name <> b1.name);
```

Scope rule: manf refers to closest nested FROM with a relation having that attribute.

The Operator ANY

- $x = \text{ANY}(\text{ <relation> })$ is a Boolean condition meaning that x equals **at least** one tuple in the relation
- Similarly, $=$ can be replaced by any of the comparison operators
 - Example: $x \geq \text{ANY}(\text{ <relation> })$ means x is not smaller than some tuples in the relation
 - Note tuples must have one component only

The Operator ALL

- $x <> \text{ALL}(\text{<relation>})$ is true if and only if for **every** tuple t in the relation, x is not equal to t
 - That is, x is not a member of the relation.
- The $<>$ can be replaced by any comparison operator
 - Example: $x \geq \text{ALL}(\text{<relation>})$ means there is no tuple larger than x in the relation

Query: From Sells(bar, beer, price), find the beer(s) sold for the highest price

```
SELECT beer  
FROM Sells  
WHERE price >= ALL(  
    SELECT price  
    FROM Sells);
```

Exercise

- From Sells(bar, beer, price), find the bar that sells the cheapest Bud Lite.

Exercise

- From Sells(bar, beer, price), find the bar that sells the cheapest Bud Lite.

Select bar

From Sells

Where beer = 'Bud Lite' and

price <= All(Select price

from Sells

where beer = 'Bud Lite')

Set operation over queries

- The default for union, intersection, and difference is **set semantics**, and is expressed by the following forms, each involving subqueries:
 - (subquery) UNION (subquery)
 - (subquery) INTERSECT (subquery)
 - (subquery) EXCEPT (subquery)
- The result is a set: duplications are removed.

Example

- **Happy Drinker:** From relations Likes(drinker, beer), Sells(bar, beer, price) and Frequent(drinker, bar), find the drinkers and beers such that:
 1. The drinker likes the beer, and
 2. The drinker frequents at least one bar that sells the beer

```
(SELECT * FROM Likes)
INTERSECT
(SELECT drinker, beer
FROM Sells, Frequent
WHERE Frequent.bar = Sells.bar
);
```

Bag semantics and set semantics

- The SELECT-FROM-WHERE statement uses **bag semantics**
 - **Selection:** preserve the number of occurrences
 - **Projection:** preserve the number of occurrences (no duplicate elimination)
 - **Cartesian product, join:** no duplicate elimination

Sells

Bar	Beer	Price
'Sue Bar'	'Bud'	2.00
'Joe Bar'	'Bud'	2.00
'Fat Daddy'	'Bud'	3.00

SELECT Price
FROM Sells



Price
2.00
2.00
3.00

Controlling Duplicate Elimination

- Force the result to be a **set** by SELECT **DISTINCT**

- From Sells(bar, beer, price), find all the different prices charged for beers:

```
SELECT DISTINCT price  
FROM Sells;
```

- Force the result to be a **bag** (i.e., don't eliminate duplicates) by **ALL**, as in ... UNION ALL ...

- Lists drinkers who frequent more bars than they like beers, and does so as many times as the difference of those counts

```
(SELECT drinker FROM Frequents)
```

```
INTERCEPT ALL
```

```
(SELECT drinker FROM Likes);
```

Aggregations

- **SUM, AVG, COUNT, MIN, and MAX** can be applied to a column in a **SELECT** clause to produce that aggregation on the column
 - e.g. COUNT(*) counts the number of tuples
- Query: From Sells(bar, beer, price), find the average price of Bud

```
SELECT AVG(price)  
FROM Sells  
WHERE beer = 'Bud'
```

Eliminating Duplicates in an Aggregation

- **DISTINCT** inside an aggregation causes duplicates to be eliminated before the aggregation
- Query: find the number of different prices charged for Bud:

```
SELECT COUNT(DISTINCT price)
```

```
FROM Sells
```

```
WHERE beer = 'Bud' ;
```

NULL's in Aggregation

- **NULL** never contributes to a sum, average, or count, and can never be the minimum or maximum of a column
 - The aggregation function ignore the rows with NULL values.
- But if there are no non-NULL values in a column, then the result of the aggregation is **NULL**

```
Select count(*)  
From Sells  
Where beer = 'Bud'
```

The number of bars that sell Bud

```
Select count(price)  
From Sells  
Where beer = 'Bud'
```

The number of bars that sell Bud at a **known** price

Group By

- We may follow a SELECT-FROM-WHERE expression by **GROUP BY** and a list of attributes
 - The relation that results from the SELECT-FROM-WHERE is grouped according to the values of all those attributes, and any aggregation is applied only within each group
- Query: From Sells(bar, beer, price), find the average price for each beer:

```
SELECT beer, AVG(price)
```

```
FROM Sells
```

```
GROUP BY beer
```

Example

- **Query:** From Sells(bar, beer, price) and Frequents (drinker, bar), find for each drinker the average price of Bud at the bars they frequent:

```
SELECT drinker, AVG(price)  
  
FROM Frequents, Sells  
  
WHERE beer = 'Bud' AND  
      Frequents.bar = Sells.bar  
  
GROUP BY drinker;
```

Compute drinker-bar-price of Bud tuples first, then group by drinker

Modifications

- A modification command does NOT return a result as a query does, but it **changes the database** in some way
- There are three kinds of modifications:
 1. **Insert** a tuple or tuples
 2. **Delete** a tuple or tuples
 3. **Update** the value(s) of an existing tuple or tuples

Insertion

- To insert a single tuple:

INSERT INTO <relation>

VALUES (<list of values>);

- Example: add to Likes(drinker, beer) the fact that Sally likes Bud:

INSERT INTO Likes

VALUES ('Sally', 'Bud');

Specifying Attributes in INSERT

- We may add to the relation **a list of attributes**
- There are two reasons to do so:
 1. We forget the standard order of attributes for the relation
 2. We don't have values for all attributes, and we want the system to fill in missing components with NULL or default values
- Another way to add the fact that Sally likes Bud to Likes(drinker, beer):

```
INSERT INTO Likes (beer, drinker)
```

```
VALUES ('Bud', 'Sally');
```

Inserting Many Tuples

- We may insert the entire result of a query into a relation, using the form:

```
INSERT INTO <relation>
( <subquery> );
```

E.g., `INSERT INTO Beers (name)`
`(SELECT beer from Sells) ;`

Example: Insert a Subquery

- Using `Frequents(drinker, bar)`, enter into the new relation **PotBuddies** (name) all of Sally's "potential buddies," i.e., those drinkers who frequent at least one bar that Sally also frequents

The other
drinker

```
INSERT INTO PotBuddies
  (SELECT d2.drinker
   FROM Frequents d1, Frequents d2
  WHERE d1.drinker = 'Sally' AND
        d2.drinker <> 'Sally' AND
        d1.bar = d2.bar
  ) ;
```

Pairs of Drinker tuples where the first is for Sally, the second is for someone else, and the bars are the same

Deletion

- To delete tuples satisfying a condition from some relation:

DELETE FROM <relation>

WHERE <condition>;

- Example: Delete from Likes(drinker, beer) the fact that Sally likes Bud:

DELETE FROM Likes

WHERE drinker = 'Sally' AND
beer = 'Bud' ;

Delete all Tuples

- Make the relation Likes empty:

```
DELETE FROM Likes;
```

- Note no WHERE clause needed

Delete Many Tuples

- Delete from Beers(name, manf) all beers for which there is another beer by the same manufacturer.

```
DELETE FROM Beers b  
WHERE EXISTS  
(  
    SELECT name FROM Beers a  
    WHERE a.manf = b.manf AND  
        a.name <> b.name  
) ;
```

Beers with the same manufacturer
and a different name from the name
of the beer represented by tuple b

Updates

- To change certain attributes in certain tuples of a relation:

UPDATE <relation>

SET <list of attribute assignments>

WHERE <condition on tuples>;

- Example: Change drinker Fred's phone number to 555-1212:

```
UPDATE Drinkers
```

```
SET phone = '555-1212'
```

```
WHERE name = 'Fred';
```

Update Several Tuples

- Increase price that is cheap:

```
UPDATE Sells
```

```
SET price = price * 1.07
```

```
WHERE price < 3.0;
```